

genotypst: A bioinformatics Typst package for sequence analysis and visualization

`genotypst` is a bioinformatics package for Typst that enables analysis and visualization of biological data. It provides functionality for parsing FASTA and Newick files and generating publication-ready visualizations, including multiple sequence alignments, sequence logos, genome maps, and phylogenetic trees.

Contents

Working with sequence data	2
Loading data	2
FASTA rendering	2
Pairwise sequence alignment	2
Rendering pairwise alignments	3
Dynamic programming matrix visualization	3
Scoring matrices	4
Multiple sequence alignments	5
Sequence logos	6
Color palettes	7
Amino acid palette	7
Nucleic acid palettes	7
Rendering genome maps and parsing GFF3 files	8
Working with phylogenetic trees	9
Visualizing trees	10
Customizing visualizations	11
Font selection	11
Specifying residue palettes	14
Bibliography	14

Working with sequence data

genotypst provides functions to parse sequence data and produce different visualizations.

Loading data

The `parse-fasta` function reads FASTA data and returns a dictionary mapping sequence identifiers to their corresponding sequences.

```
#let sequences = parse-fasta(read("/docs/data/dna.fna"))
```

```
(  
  seq_1: "AAGGGACACTGATTTTCTCCACAGCTGGGCCGTGGACCGTAGTGTTTCAAGAAGCCCACACAC",  
  seq_2: "GCAATGGAGACAACATAGCCAACTACCTACTAGATGGCCTAGATCTGCCGCA",  
  seq_3: "GGAAGTGGCGTTACAGACAGTTGTGAGCCACCACATGGGCCTGGGATTTAAATATTATAAAGCTCCTC",  
)
```

FASTA rendering

Use `render-fasta` to display sequences in the standard FASTA format.

```
#context {  
  set text(size: 0.8em)  
  render-fasta(sequences, max-width: 50)  
}
```

```
>seq_1  
AAGGGACACTGATTTTCTCCACAGCTGGGCCGTGGACCGTAGTGTTTCA  
AGAAGCCCACACAC  
>seq_2  
GCAATGGAGACAACATAGCCAACTACCTACTAGATGGCCTAGATCTGCCG  
CA  
>seq_3  
GGAAGTGGCGTTACAGACAGTTGTGAGCCACCACATGGGCCTGGGATTTA  
AATATTATAAAGCTCCTC
```

In this example, `max-width` controls how many characters appear per line (default is 60).

Pairwise sequence alignment

Pairwise alignment is a method for comparing biological sequences, allowing quantification of sequence similarity and identification of evolutionary relationships between the residues of two sequences. `genotypst` supports both global alignment (end-to-end)¹ and local alignment (best-matching subsequences)², using dynamic programming with a user-defined scoring scheme (match/mismatch or a substitution matrix) and gap penalties.

The `align-seq-pair` function performs pairwise alignment using either match/mismatch scores or a substitution matrix (see the scoring matrices section below).

```
#let dna_pair_alignment = align-seq-pair(
  "ACT",
  "ACGT",
  match-score: 3,
  mismatch-score: -1,
  gap-penalty: -2,
  mode: "global",
)

#let protein_pair_alignment = align-seq-pair(
  "MAVHLTPPEKS",
  "HLTPEE",
  scoring-matrix: "BLOSUM62",
  gap-penalty: -4,
  mode: "local",
)
```

The DNA example uses a custom match/mismatch model for a global alignment, while the protein example uses BLOSUM62 for a local alignment. The alignment scores are 7 and 28.

Rendering pairwise alignments

Pairwise alignments can be rendered using the `render-pair-alignment` function. The example below uses the local protein alignment from the previous section, showing the unaligned regions using a light gray color.

```
// Because there may be multiple optimal alignments, `traceback-paths`
// is an array. We use `.at(0)` to visualize the path of the first alignment.
#render-pair-alignment(
  protein_pair_alignment.seq-1,
  protein_pair_alignment.seq-2,
  protein_pair_alignment.traceback-paths.at(0),
  unaligned-color: luma(75%),
)
```

```

MGRHMTYPEEKS
 |  |  |  |
H L T - P E E

```

Local protein alignment with unaligned regions shown in light gray.

Dynamic programming matrix visualization

Dynamic programming is the core procedure used by the pairwise alignment algorithm: it fills a matrix of optimal scores for all prefix pairs of the two sequences, where each cell stores the best score achievable at that position and arrows indicate the traceback directions that can lead to an optimal alignment. The `render-dp-matrix` function renders the DP matrix of a given alignment, overlaying the traceback path used to produce the final alignment.

```
// Pass traceback path directly from align-seq-pair output (end-to-start).
#render-dp-matrix(
  dna_pair_alignment.seq-1,
  dna_pair_alignment.seq-2,
  cell-values: dna_pair_alignment.dp-matrix.scores,
  path: dna_pair_alignment.traceback-paths.at(0),
  arrows: dna_pair_alignment.dp-matrix.arrows,
)
```

	-	A	C	G	T
-	0	-2	-4	-6	-8
A	-2	3	1	-1	-3
C	-4	1	6	4	2
T	-6	-1	4	5	7

Dynamic programming matrix for the DNA alignment, with the optimal path highlighted.

Scoring matrices

Substitution matrices assign scores for aligning residues in pairwise and multiple sequence alignments, rewarding likely substitutions and penalizing unlikely ones. They represent substitution preferences as log-odds scores derived from observed evolutionary changes. In this way, they guide alignment algorithms by quantifying how plausible it is for one residue to replace another over time, helping produce biologically meaningful alignments. `genotypst` provides scoring matrices from the BLOSUM³ and PAM⁴ families for protein sequences, as well as the EDNAFULL matrix for DNA and RNA sequences.

Use `get-scoring-matrix` to retrieve a matrix data and `get-score-from-matrix` to query scores for specific residue pairs.

```
#let blosum62 = get-scoring-matrix("BLOSUM62")
#let ar_score = get-score-from-matrix(blosum62, "L", "D")
#text[The substitution score for L vs D in BLOSUM62 is: #ar_score.]
```

The substitution score for L vs D in BLOSUM62 is: -4.

You can render the entire scoring matrix using `render-scoring-matrix` function. In the example below `scale-limit: 7` set the color scale limits to -7 to 7.

```
#context {
  set text(size: 0.9em)
  render-scoring-matrix(
    blosum62,
    scale-limit: 7,
  )
}
```

	A	R	N	D	C	Q	E	G	H	I	L	K	M	F	P	S	T	W	Y	V
A	4																			
R	-1	5																		
N	-2	0	6																	
D	-2	-2	1	6																
C	0	-3	-3	-3	9															
Q	-1	1	0	0	-3	5														
E	-1	0	0	2	-4	2	5													
G	0	-2	0	-1	-3	-2	-2	6												
H	-2	0	1	-1	-3	0	0	-2	8											
I	-1	-3	-3	-3	-1	-3	-3	-4	-3	4										
L	-1	-2	-3	-4	-1	-2	-3	-4	-3	2	4									
K	-1	2	0	-1	-3	1	1	-2	-1	-3	-2	5								
M	-1	-1	-2	-3	-1	0	-2	-3	-2	1	2	-1	5							
F	-2	-3	-3	-3	-2	-3	-3	-3	-1	0	0	-3	0	6						
P	-1	-2	-2	-1	-3	-1	-1	-2	-2	-3	-3	-1	-2	-4	7					
S	1	-1	1	0	-1	0	0	0	-1	-2	-2	0	-1	-2	-1	4				
T	0	-1	0	-1	-1	-1	-1	-2	-2	-1	-1	-1	-1	-2	-1	1	5			
W	-3	-3	-4	-4	-2	-2	-3	-2	-2	-3	-2	-3	-1	1	-4	-3	-2	11		
Y	-2	-2	-2	-3	-2	-1	-2	-3	2	-1	-1	-2	-1	3	-3	-2	-2	2	7	
V	0	-3	-3	-3	-1	-2	-2	-3	-3	3	1	-2	1	-1	-2	-2	0	-3	-1	4

BLOSUM62 scoring matrix

Multiple sequence alignments

The `render-msa` function displays multiple sequence alignments with optional residue coloring and conservation bars.

In the example below:

- `colors: true` enables residue coloring based on biochemical properties.
- `show-conservation: true` adds bars above the alignment columns to indicate the residue conservation at each position of the alignment.
- `start: 100` and `end: 145` limit the display to a specific region of interest (residues 100 to 145).

```
#let protein_msa = parse-fasta(read("/docs/data/msa.afa"))

#context {
  set text(size: 0.8em)
  render-msa(
    protein_msa,
    start: 100,
    end: 145,
    colors: true,
    show-conservation: true,
  )
}
```



MSA visualization for positions 100–145, with residue coloring and bars indicating indicating residue conservation.

Residue coloring represents amino acid physicochemical properties. The sequence alphabet (amino acid, DNA, or RNA) is determined automatically and a suitable color palette is applied.


The bars above the alignment indicate the degree of residue conservation at each column.

Sequence logos

Sequence logos⁵ are graphical summaries of residue variation and conservation across positions in a sequence alignment and are commonly used to visualize binding sites, motifs, and functional domains. In a sequence logo, the total height of each stack represents the information content (in bits) at that position, while individual letters are scaled according to their relative frequencies.

In the example below, we visualize the same region as the MSA of the previous section.

```
#render-sequence-logo(protein_msa, start: 100, end: 145)
```



Sequence logo for positions 100–145, showing conservation and residue frequency.

Like `render-msa`, `render-sequence-logo` automatically applies the appropriate color palette based on the sequence alphabet.

Color palettes

`genotypst` uses predefined color palettes to assign colors to sequence residues.

Amino acid palette

Amino acids are colored according to their physicochemical properties. Grouping residues by color helps reveal the chemical nature of conserved positions (e.g., whether a position is consistently hydrophobic or charged), which is often important for understanding protein structure, function, and evolution.

Hydrophobic

Alanine	Ala	A	#4d78ff
Histidine	His	H	#4d78ff
Isoleucine	Ile	I	#4d78ff
Leucine	Leu	L	#4d78ff
Methionine	Met	M	#4d78ff
Valine	Val	V	#4d78ff

Polar

Serine	Ser	S	#00c990
Threonine	Thr	T	#00c990
Glutamine	Gln	Q	#00c990
Asparagine	Asn	N	#00c990

Aromatic

Phenylalanine	Phe	F	#bac1d2
Tryptophan	Trp	W	#bac1d2
Tyrosine	Tyr	Y	#bac1d2

Negatively charged

Aspartic acid	Asp	D	#ff07b8
Glutamic acid	Glu	E	#ff07b8

Positively charged

Lysine	Lys	K	#e44356
Arginine	Arg	R	#e44356

Cysteine

Cysteine	Cys	C	#494e5b
----------	-----	---	---------

Glycine

Glycine	Gly	G	#f59116
---------	-----	---	---------

Proline

Proline	Pro	P	#d9de09
---------	-----	---	---------

Nucleic acid palettes

The DNA and RNA palettes assign a distinct color to each nucleotide.

DNA palette

Adenine	A	#00c990
Cytosine	C	#4d78ff
Guanine	G	#ff07b8
Thymine	T	#f59116

RNA palette


Adenine	A	#00c990
Cytosine	C	#4d78ff
Guanine	G	#ff07b8
Uracil	U	#494e5b

Rendering genome maps and parsing GFF3 files

Genome maps enable visualization of the genes and other genomic elements within a locus, highlighting their order, orientation, and length. `genotypst` provides a `render-genome-map` function that produces a genome map from an array of dictionaries, each representing a genomic feature that will be plotted:

- `start` (required): Start coordinate (1-indexed integer).
- `end` (required): End coordinate (1-indexed integer).
- `strand`: Feature orientation (1 or "+" for the positive strand, -1 or "-" for the negative strand). `none` draws an undirected block.
- `label`: Feature label
- `color`: Fill color.

```
#let f_plasmid_locus = (  
  (start: 65556, end: 66065, strand: -1, label: [_ygfA_]),  
  (start: 66118, end: 66407, label: [_oriT_], color: rgb("#696975")),  
  (start: 66479, end: 66862, strand: 1, label: [_traM_], color: rgb("#62B9F2")),  
  (start: 66977, end: 67055, strand: -1, label: [_finP_]),  
  (start: 67049, end: 67738, strand: 1, label: [_traJ_], color: rgb("#F7ED6C")),  
  (start: 67837, end: 68232, strand: 1, label: [_traY_]),  
  (start: 68265, end: 68630, strand: 1, label: [_traA_]),  
  (start: 68645, end: 68956, strand: 1, label: [_traL_]),  
  (start: 68978, end: 69544, strand: 1, label: [_traE_]),  
)  
  
#render-genome-map(  
  f_plasmid_locus,  
  coordinate-axis: true,  
  unit: "bp",  
)
```




If your genomic features are stored in a GFF3 file, you can use the `parse-gff` function to parse the file and return an array of dictionaries suitable as input to `render-genome-map`.

In the example below:

- `feature-types`: Return only features of the specified types.
- `range`: Return only features within the specified genomic range, defined as (accession, start, end).
- `exclude-partial`: Exclude genes that are not fully contained within the specified range.
- `label-attribute`: Use this feature attribute as the label in the genome map.

```
#let parsed_gff3_data = parse-gff(  
  read("data/NC_002483.gff3"),  
  feature-types: ("CDS", "pseudogene"),  
  range: ("NC_002483.1", 68960, 82550),  
  exclude-partial: true,  
  label-attribute: "product",  
)  
  
#render-genome-map(  
  parsed_gff3_data,  
  coordinate-axis: true,  
  scale-bar: true,  
  unit: "bp",  
)
```



Working with phylogenetic trees

`genotypst` includes functions to parse and render phylogenetic trees. Trees can be created by parsing Newick-formatted strings with `parse-newick` or by manually constructing nested dictionary structures.

```
#let parsed_newick_tree = parse-newick(
  "((('Leaf A':0.2,'Leaf B':0.1)'Internal node':0.3,'Leaf C':0.6)Root;"
)

#let manual_tree = (
  rooted: true,
  name: text(weight: "bold")[Root],
  length: none,
  children: (
    (
      name: text(fill: orange)[Internal node],
      length: 0.3,
      children: (
        (name: text[_Leaf A_], length: 0.2, children: none),
        (name: "Leaf B", length: 0.1, children: none),
      ),
    ),
    (name: text(fill:blue)[Leaf C], length: 0.6, children: none),
  ),
)
```

Visualizing trees



genotypst can produce visualizations of phylogenetic trees. To illustrate this, we will read and render a Newick file containing a phylogeny of the *Hominoidea* superfamily, which was extracted from the Ensembl Compara species tree⁶.

```
#let hominoidea_tree = parse-newick(read("/docs/data/hominoidea.nwk"))
```

Two types of trees can be produced: rectangular trees (`render-rectangular-tree`) and unrooted trees (`render-unrooted-tree`). In the examples below:

- `tip-label-italics: true` renders tip labels in italics.
- `orientation: "horizontal"` sets the rectangular tree to be oriented horizontally. It can also be set to `"vertical"` for a vertical layout.
- `align-tip-labels: true` aligns tip labels and connects them to branches with dotted leader lines.
- `hide-internal-labels: true` hides internal node labels.
- `layout: "daylight"` applies the daylight layout algorithm to the unrooted tree. The default layout is `"equal-angle"`.

```
#context {
  set text(size: 0.9em)
  render-rectangular-tree(
    hominoidea_tree,
    tip-label-italics: true,
    orientation: "horizontal",
    align-tip-labels: true,
  )
  render-unrooted-tree(
    hominoidea_tree,
    tip-label-italics: true,
    hide-internal-labels: true,
    layout: "daylight",
  )
}
```



Customizing visualizations

Font selection

By default, the visualizations produced by `genotypst` are rendered using the default document font. To control font size or font family for these visualizations, wrap the rendering function in a `context` block and set the desired text properties.

```
#let dna_msa = (
  "seq1": "AGTCTCAAGATAACTTTGAAACAAAC",
  "seq2": "AGTTTCCAAGTGGATTTGGAATTGAAC",
  "seq3": "ACTCT-CGGATGGATTCCGATACAAAC",
  "seq4": "AGTCT---GATTGATGTGGATACAAAC",
  "seq5": "AGTCT--GGGTGGATTTGG-AACAAAT",
  "seq6": "CAGTGCTCCCTGGTGGTGG-ACCATCT",
  "seq7": "AGTCTCAAGACGGATACTG--ATGCCC",
)

#context {
  set text(font: "Maple Mono", size: 0.785em)
  render-msa(dna_msa)
}
```

```
seq1 AGTCTCAAGATAACTTTGAAACAAAC
seq2 AGTTTCCAAGTGGATTTGGAATTGAAC
seq3 ACTCT-CGGATGGATTCCGATACAAAC
seq4 AGTCT- - -GATTGATGTGGATACAAAC
seq5 AGTCT- - GGGTGGATTTGG-AACAAAT
seq6 CAGTGCTCCCTGGTGGTGG-ACCATCT
seq7 AGTCTCAAGACGGATACTG- -ATGCCC
```

Default document font.

```
seq1 AGTCTCAAGATAACTTTGAAACAAAC
seq2 AGTTTCCAAGTGGATTTGGAATTGAAC
seq3 ACTCT-CGGATGGATTCCGATACAAAC
seq4 AGTCT---GATTGATGTGGATACAAAC
seq5 AGTCT--GGGTGGATTTGG-AACAAAT
seq6 CAGTGCTCCCTGGTGGTGG-ACCATCT
seq7 AGTCTCAAGACGGATACTG--ATGCCC
```

Custom font (Maple Mono).

```
#context {
  show text: set text(font: "Libertinus Serif")
  render-pair-alignment(
    protein_pair_alignment.seq-1,
    protein_pair_alignment.seq-2,
    protein_pair_alignment.traceback-paths.at(0),
  )
}
```

```
MGRHMTYPEEKS
 | | |||
HLT-PEE
```

Default document font.

```
MGRHMTYPEEKS
 | | |||
HLT-PEE
```


Custom font (Libertinus Serif).

```
#context {
  show text: set text(font: "Libertinus Serif")
  render-sequence-logo(dna_msa)
}
```


Default document font.

Custom font (Libertinus Serif).

```
#context {
  show text: set text(font: "Libertinus Serif")
  render-genome-map(f_plasmid_locus)
}
```




Default document font.




Custom font (Libertinus Serif).

```
#let hominoidea_tree = parse-newick(read("/docs/data/hominoidea.nwk"))
#context {
  show text: set text(font: "Libertinus Serif", size: 0.9em)
  #render-rectangular-tree(hominoidea_tree, tip-label-italics: true)
}
```



Default document font.



Custom font (Libertinus Serif).

Specifying residue palettes


If you want to use an alternative palette to color residues, you can provide a dictionary to the `palette` parameter of `render-msa` and `render-sequence-logo`. The provided palette can be either custom or one provided by `genotypst` under `residue-palette`.

The following palettes are available:

- **Protein:** default (8 colors), dayhoff (6), zappo (7), takabatake4 (4), takabatake5 (5), takabatake6 (6), takabatake7 (7), and takabatake8 (8)⁷.
- **DNA/RNA:** default (4), gc (2), purine-pyrimidine (2)

For example, to use the Dayhoff amino acid color palette in a sequence logo:

```
#render-sequence-logo(  
  protein_msa,  
  start: 100,  
  end: 145,  
  palette: residue-palette.aa.dayhoff,  
)
```



Sequence logo for positions 100–145 using the Dayhoff amino acid color palette.

Bibliography

1. Needleman, S. B. & Wunsch, C. D. A general method applicable to the search for similarities in the amino acid sequence of two proteins. *Journal of Molecular Biology* **48**, 443–453 (1970).
2. Smith, T. & Waterman, M. Identification of common molecular subsequences. *Journal of Molecular Biology* **147**, 195–197 (1981).
3. Henikoff, S. & Henikoff, J. G. Amino acid substitution matrices from protein blocks. *Proceedings of the National Academy of Sciences* **89**, 10915–10919 (1992).
4. Dayhoff, M., Schwartz, R. & Orcutt, B. A model of evolutionary change in proteins. in *Atlas of protein sequence and structure* (National Biomedical Research Foundation, Washington, D.C, 1979).
5. Schneider, T. D. & Stephens, R. Sequence logos: a new way to display consensus sequences. *Nucleic Acids Research* **18**, 6097–6100 (1990).
6. Herrero, J. *et al.* Ensembl comparative genomics resources. *Database* **2016**, bav96 (2016).
7. Takabatake, K. *et al.* Improved Large-Scale Homology Search by Two-Step Seed Search Using Multiple Reduced Amino Acid Alphabets. *Genes* **12**, 1455 (2021).